

CITS3402/CITS5507 Assignment 1: Parallel Hash Collision Attack

[Your Name], 24914408

Semester 2, 2026

Birthday-attack algorithm and collision-detection data structure

`toy_hash` produces a 48-bit output, so by the birthday bound a matching pair between two independent families of nonce trials is expected after roughly $\sqrt{\pi/2 * 2^{48}} \sim 2.1e7$ combined trials — far fewer than 2^{48} for a preimage search. Rather than fixing one file’s nonce and searching only the other, the program grows two pools of trials in parallel, one per file, stopping once a trial’s hash already exists in the *other* file’s pool.

Both pools live in one shared open-addressing hash table keyed by the 48-bit hash. Each slot stores the hash, the nonce that produced it, which file it came from, and a “ready” flag. Insertion mixes the hash with a Fibonacci multiplier, masks it to the (power-of-two) table size to pick a home slot, then linear probes to the first free or matching slot. A match already owned by the *other* file is a genuine collision; one owned by the *same* file is a redundant self-collision, discarded. Capacity is four times the trial budget ($4 * 2^{26}$ slots by default) so load factor stays at or below 0.5 and probe chains stay short. Nonces are generated sequentially via an atomic counter per file rather than randomly — `toy_hash`’s MurmurHash3-style finalising mix already scatters sequential inputs uniformly, giving the same statistics as random sampling while wasting no nonce on a repeat.

Parallelisation and synchronisation with OpenMP

Threads split evenly into an “A-group” and a “B-group” (even/odd thread id) inside one `#pragma omp parallel` region; each hashes trials only for its file, drawing the next nonce with `atomic_fetch_add` on a per-file counter — no two threads ever hash the same nonce, and none blocks.

The shared hash table is the only structure threads communicate through, made safe without any `critical` section or `omp_lock_t`. Each slot’s hash field is a C11 `_Atomic uint64_t` written only via `atomic_compare_exchange_strong` from `TABLE_EMPTY` to the trial’s hash; exactly one thread can ever win that CAS for a given slot, so that slot’s `nonce/owner` fields are subsequently written by that thread alone — no two threads ever write the same memory. The remaining hazard is a reader observing a slot as claimed before the winner finishes writing `nonce/owner`; this is closed with a “published” flag — the winner writes `nonce/owner`, then stores `ready = 1` with `memory_order_release`, and a reader spins on `ready` with `memory_order_acquire` before reading them, so it never observes a torn write. That spin window is only the few instructions between the CAS succeeding and the stores after it, so it costs nothing in practice. A second atomic (`found_flag`) is claimed with one final CAS so only the first thread to find a collision records the winning nonce pair; every other thread observes the shared `stop_flag` and exits. Before writing

any file, the program independently recomputes `toy_hash` on the final bytes and aborts if they disagree, so a bug in the concurrent path can never silently produce an invalid submission.

Memory requirements and trade-offs

The dominant cost is the collision table: at the default 2^{26} -trials-per-side budget, capacity is 2^{28} slots, each holding an 8-byte hash, 8-byte nonce, and two 1-byte tags — about 4.3 GB. This beats a smaller, higher-load-factor table because linear probing degrades sharply as load factor approaches 1, and Kaya nodes have far more memory than this needs, so speed is favoured over economy. A second, much smaller cost is per-thread: each thread keeps one private, mutable copy of whichever file(s) it hashes (loaded once, header patched every trial), so only the 16-byte nonce is rewritten per trial rather than re-copying the whole file — for the largest (~900 KB) pair at 96 threads this is under 90 MB, negligible next to the table.

Performance metrics and analysis

`search_seconds` on a 10-core development machine (every result independently verified against `check_toy_hash.py`), for the smallest pair, confirms both correctness and near-linear early scaling:

threads	search_seconds	speedup
1	21.53	1.00x
2	10.69	2.02x
4	6.00	3.59x
8	3.78	5.69x
10	3.42	6.29x

Both runs found ~10.39M trials per side (~20.8M combined), matching the birthday bound almost exactly. Efficiency drops from 100% at 2 threads to ~63% at 10: the search is memory-bandwidth-bound (every insert/probe is an effectively random access into a multi-gigabyte table), so once all cores are active they contend for cache/DRAM bandwidth rather than compute — a trend expected to continue at Kaya’s higher core counts.

Solving the full pairs end-to-end (student ID 24914408, verified against `check_toy_hash.py`) shows the other source of variance: how many trials the *specific* target hash needs. `1_kilo` (64 KB) took 999.97 s at 25.32M trials/side, while the larger `2_mega` (128 KB) took only 747.5 s at 11.41M trials/side — despite costing roughly twice as much per hash, it needed under half the trials. Trial count has a long right tail around its ~2.1e7 combined mean, so wall-clock time is not a deterministic function of file size alone and must be measured per pair. This is also why harder pairs need Kaya’s full 96-core node: `3_giga-6_exa` cost several times more per hash, so fitting the 900 s budget even on an unlucky run needs that throughput.

[Insert Kaya results here: `search_seconds` for all six pairs across thread counts 1–96 from `scripts/scaling_job.slurm`, with the resulting speedup/efficiency curves.]